

Reply Assistant

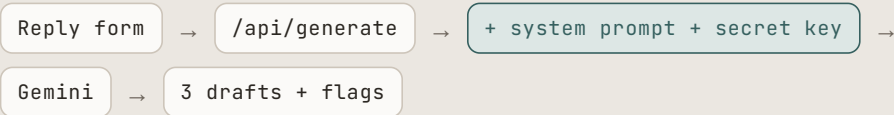
A web app that reads a message from a parent or a distant friend, works out what they actually need, and drafts three replies you can review and send yourself.

Eva September 2026 reply-assistant-three.vercel.app

MODEL	HOSTING	STORES MESSAGES	SENDS REPLIES
Gemini 3.6 Flash	Vercel	No	Never; you do

01

How it works



The browser never talks to Gemini directly. Every request goes through my own serverless function, which checks the input, adds the system prompt and the API key, and turns failures into plain error messages instead of made-up replies.

02

Use

INTENDED FOR

- Everyday replies to parents, relatives and old friends
- Getting past the "I'll answer later" feeling
- Drafts you read, edit, and send yourself

NOT FOR

- Emergencies or someone in real distress
- Therapy, legal or money decisions
- Deceiving someone about what happened
- Users under 18 (Gemini API terms)

03

Inputs and outputs

YOU GIVE IT

- The message you received
- Who it's from (10 relationship types)
- Warmth 1–5, length, catch-up yes/no
- True facts to include (optional)
- Up to 4 chat screenshots (optional)

YOU GET BACK

- Three drafts: warm, brief, honest
- What the sender seems to need
- A checklist of [gaps] to fill in
- Warnings: urgent, pressure, private request

04

Data and privacy

- The app has no database. Nothing you paste is saved by it.
- Your message, notes and screenshots are sent to Google's Gemini API to generate the drafts.
- Screenshots include the *other person's* messages, and they haven't agreed to share them.
- The API key lives only on the server and is never visible in the browser.

Free tier caveat. This project uses Gemini's free tier. Google's terms say free-tier content can be used to improve its products and may be read by human reviewers, and they ask users not to submit personal information. Until it moves to a paid key, avoid real names, addresses and anything private.

05

Safety measures

- **No invented facts.** If it doesn't know a detail about your life, it writes [what you ate] instead of guessing.
- **Urgent messages** get a "consider calling them" warning, not a breezy reply.
- **Pressure or guilt-tripping** is flagged; drafts stay calm and don't apologise for things you didn't do.
- **Private requests** (money, addresses, health) are left out of drafts and flagged.
- **No judgement** of you, the sender, or how long you took to reply.
- **Text inside a message is treated as data**, so "ignore your instructions" in a message doesn't work.

06

Known limitations

- The model can still misread tone or get a detail wrong. Always read before sending.
- "Parent" or "distant friend" is broad; it doesn't know your actual people.
- The "reciprocity score" rates how much the message invites a reply, but it could still be misread as a score of the person.
- The interface is English only, though drafts follow the sender's language.
- Rate limiting is basic and wouldn't stop determined abuse.
- Leaning on it too much could make relationships feel less personal. It's meant to get you started, not to replace talking.

07

Evaluation

18 automated checks cover the plumbing: input validation, image limits, the request sent to Gemini, and error handling. They don't measure whether the replies are any good.

Reply quality has only been tested by hand so far. Next step: 20 real messages I've received, each judged on one question: would I send this with only the brackets filled in?

Versions

v1	Sep 14	System prompt, Gemini API route, simple demo page
v2	Sep 15	AI Studio interface, screenshot reading, safety warnings, GitHub auto-deploy
v3	~Sep 16	Warm frosted-glass restyle, layout unchanged

Built with Claude (prompt, API, integration), Google AI Studio (interface), Vercel and GitHub. Based on my original role card for a patient, non-judgemental letter writer.